

Exp6: Generative Adversarial Networks

使用生成对抗网络(GAN)生成 图片

```
In [1]: 1 from __future__ import print_function
2
3 # import argparse
4 import os
5 import random
6 import torch
7 import torch.nn as nn
8 # import torch.nn.parallel
9 # import torch.backends.cudnn as cudnn
10 import torch.optim as optim
11 import torch.utils.data
12 import torchvision.datasets as dset
13 import torchvision.transforms as transforms
14 import torchvision.utils as vutils
15 import numpy as np
16 import matplotlib.pyplot as plt
17 import matplotlib.animation as animation
18 from IPython.display import HTML
19 %matplotlib inline
20 %config InlineBackend.figure_format = 'svg'
21
22 Seed = 220103
23 random.seed(Seed)
24 torch.manual_seed(Seed)
```

executed in 1.78s, finished 22:42:17 2022-01-05

Out[1]: <torch._C.Generator at 0x1eaa2a16730>

1. 数据准备

- 将所有人脸图片放入 face 文件夹 (共13233)
- 创建数据迭代器

```

In [2]: 1 dataroot = r"F:\2021研一上学期\深度学习\assignment6-GAN\image_data"
2 image_size = 64
3 batch_size = 32
4
5 # Create the dataset
6 dataset = dset.ImageFolder(root=dataroot,
7                             transform=transforms.Compose([
8                                 transforms.Resize(image_size),
9                                 transforms.CenterCrop(image_size),
10                                transforms.ToTensor(),
11                                transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
12                                ]))
13 # Create the dataloader
14 dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
15                                           shuffle=True, num_workers=2)
16
17 # Decide which device we want to run on
18 device = torch.device("cuda:0" if (torch.cuda.is_available()) else "cpu")
19
20 # Plot some training images
21 real_batch = next(iter(dataloader))
22 plt.figure(figsize=(8,8))
23 plt.axis("off")
24 plt.title("Training Images")
25 plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=2, normalize=True).cpu(),
executed in 6.15s, finished 22:42:23 2022-01-05

```

Out[2]: <matplotlib.image.AxesImage at 0x1eaabfbb490>



```

In [3]: 1 # Number of channels in the training images. For color images this is 3
2 nc = 3
3
4 # Size of z latent vector (i.e. size of generator input)
5 nz = 300
6
7 # Size of feature maps in generator
8 ngf = 64
9
10 # Size of feature maps in discriminator
11 ndf = 64
12
13 # Number of training epochs
14 num_epochs = 300
15
16 # Learning rate for optimizers
17 lr = 0.0002

```

executed in 14ms, finished 22:42:23 2022-01-05

2. 模型搭建

在生成网络G和对抗网络D上调用自定义权重初始化


在DCGAN论文中，作者指定所有模型的权值都应该从均值为0,标准差0.02的正态分布随机初始化。

weights_init函数接受一个初始化的模型作为输入，并重新初始化所有卷积、卷积转置和批处理规范化层，以满足这个标准。此函数在初始化后立即应用于模型。

```
In [4]: 1 # custom weights initialization called on netG and netD
2 def weights_init(m):
3     classname = m.__class__.__name__
4     if classname.find('Conv') != -1:
5         nn.init.normal_(m.weight.data, 0.0, 0.02)
6     elif classname.find('BatchNorm') != -1:
7         nn.init.normal_(m.weight.data, 1.0, 0.02)
8         nn.init.constant_(m.bias.data, 0)
```

executed in 14ms, finished 22:42:23 2022-01-05

生成器 G 的结构如下



其中，nz为输入向量z的长度，ngf与通过生成器传播的特征图的大小有关，nc为输出图像中的通道数(对于RGB图像，设置为3)。

```
In [5]: 1 # Generator Code
2
3 class Generator(nn.Module):
4     def __init__(self):
5         super(Generator, self).__init__()
6         self.main = nn.Sequential(
7             # input is Z, going into a convolution
8             nn.ConvTranspose2d( nz, ngf * 8, 4, 1, 0, bias=False),
9             nn.BatchNorm2d(ngf * 8),
10            nn.ReLU(True),
11            # state size. (ngf*8) x 4 x 4
12            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
13            nn.BatchNorm2d(ngf * 4),
14            nn.ReLU(True),
15            # state size. (ngf*4) x 8 x 8
16            nn.ConvTranspose2d( ngf * 4, ngf * 2, 4, 2, 1, bias=False),
17            nn.BatchNorm2d(ngf * 2),
18            nn.ReLU(True),
19            # state size. (ngf*2) x 16 x 16
20            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
21            nn.BatchNorm2d(ngf),
22            nn.ReLU(True),
23            # state size. (ngf) x 32 x 32
24            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
25            nn.Tanh()
26            # state size. (nc) x 64 x 64
27        )
28
29    def forward(self, input):
30        return self.main(input)
```

executed in 14ms, finished 22:42:23 2022-01-05

实例化生成器并应用weights_init函数

```
In [6]: 1 # Create the generator
2 netG = Generator().to(device)
3
4 # Apply the weights_init function to randomly initialize all weights
5 # to mean=0, stdev=0.02.
6 netG.apply(weights_init)
7
8 # Print the model
9 print(netG)
```

executed in 45ms, finished 22:42:23 2022-01-05

```
Generator(
  (main): Sequential(
    (0): ConvTranspose2d(300, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (11): ReLU(inplace=True)
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (13): Tanh()
  )
)
```

鉴定器 D 是一个二值分类网络，它以图像作为输入，并输出输入图像是真实的(而不是假的)的概率。

在这里，D 获取一个3x64x64的输入图像，通过一系列Conv2d、BatchNorm2d和LeakyReLU层处理它，并通过Sigmoid激活函数输出最终的概率。

```
In [7]: 1 class Discriminator(nn.Module):
2     def __init__(self):
3         super(Discriminator, self).__init__()
4         self.main = nn.Sequential(
5             # input is (nc) x 64 x 64
6             nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
7             nn.LeakyReLU(0.2, inplace=True),
8             # state size. (ndf) x 32 x 32
9             nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
10            nn.BatchNorm2d(ndf * 2),
11            nn.LeakyReLU(0.2, inplace=True),
12            # state size. (ndf*2) x 16 x 16
13            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
14            nn.BatchNorm2d(ndf * 4),
15            nn.LeakyReLU(0.2, inplace=True),
16            # state size. (ndf*4) x 8 x 8
17            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
18            nn.BatchNorm2d(ndf * 8),
19            nn.LeakyReLU(0.2, inplace=True),
20            # state size. (ndf*8) x 4 x 4
21            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
22            nn.Sigmoid()
23        )
24
25    def forward(self, input):
26        return self.main(input)
```

executed in 13ms, finished 22:42:23 2022-01-05

```
In [8]: 1 # Create the Discriminator
2 netD = Discriminator().to(device)
3
4 # Apply the weights_init function to randomly initialize all weights
5 # to mean=0, stdev=0.2.
6 netD.apply(weights_init)
7
8 # Print the model
9 print(netD)
```

executed in 29ms, finished 22:42:23 2022-01-05

```
Discriminator(
  (main): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): LeakyReLU(negative_slope=0.2, inplace=True)
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (12): Sigmoid()
  )
)
```

接下来，我们将真正的标签定义为1，假标签定义为0。在计算生成器和鉴定器的loss时会用到这些标签，这也是原GAN论文中使用的约定。

最后，设置了两个单独的优化器，一个用于D，一个用于G，如DCGAN论文所述，都是Adam优化器，学习速率为0.0002, $\beta_1 = 0.5$

为了跟踪生成器的学习进程，从高斯分布(即fixed_noise)中生成一批固定的潜在向量。

```
In [9]: 1 # Initialize BCELoss function
2 criterion = nn.BCELoss()
3
4 # Create batch of latent vectors that we will use to visualize
5 # the progression of the generator
6 fixed_noise = torch.randn(64, nz, 1, 1, device=device)
7
8 # Establish convention for real and fake labels during training
9 real_label = 1.
10 fake_label = 0.
11
12 # Setup Adam optimizers for both G and D
13 beta1 = 0.5
14 optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
15 optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))
```

executed in 14ms, finished 22:42:23 2022-01-05

3. 模型训练

```

In [10]: 1 # Training Loop
2
3 # Lists to keep track of progress
4 img_list = []
5 fake_list = []
6 G_losses = []
7 D_losses = []
8 iters = 0
9
10 print("Starting Training Loop...")
11 # For each epoch
12 for epoch in range(num_epochs):
13     # For each batch in the dataloader
14     for i, data in enumerate(dataloader, 0):
15
16         #####
17         # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
18         #####
19         ## Train with all-real batch
20         netD.zero_grad()
21         # Format batch
22         real_cpu = data[0].to(device)
23         b_size = real_cpu.size(0)
24         label = torch.full((b_size,), real_label, dtype=torch.float, device=device)
25         # Forward pass real batch through D
26         output = netD(real_cpu).view(-1)
27         # Calculate loss on all-real batch
28         errD_real = criterion(output, label)
29         # Calculate gradients for D in backward pass
30         errD_real.backward()
31         D_x = output.mean().item()
32
33         ## Train with all-fake batch
34         # Generate batch of latent vectors
35         noise = torch.randn(b_size, nz, 1, 1, device=device)
36         # Generate fake image batch with G
37         fake = netG(noise)
38         label.fill_(fake_label)
39         # Classify all fake batch with D
40         output = netD(fake.detach()).view(-1)
41         # Calculate D's loss on the all-fake batch
42         errD_fake = criterion(output, label)
43         # Calculate the gradients for this batch, accumulated (summed) with previous gradients
44         errD_fake.backward()
45         D_G_z1 = output.mean().item()
46         # Compute error of D as sum over the fake and the real batches
47         errD = errD_real + errD_fake
48         # Update D
49         optimizerD.step()
50
51         #####
52         # (2) Update G network: maximize log(D(G(z)))
53         #####
54         netG.zero_grad()
55         label.fill_(real_label) # fake labels are real for generator cost
56         # Since we just updated D, perform another forward pass of all-fake batch through D
57         output = netD(fake).view(-1)
58         # Calculate G's loss based on this output
59         errG = criterion(output, label)
60         # Calculate gradients for G
61         errG.backward()
62         D_G_z2 = output.mean().item()
63         # Update G
64         optimizerG.step()
65
66         # Output training stats
67         if i % 50 == 0:
68             print(' [%d/%d] [%d/%d] \tLoss_D: %.4f \tLoss_G: %.4f \tD(x): %.4f \tD(G(z)): %.4f / %.4f'
69                   % (epoch, num_epochs, i, len(dataloader),
70                     errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))
71
72         # Save Losses for plotting later
73         G_losses.append(errG.item())
74         D_losses.append(errD.item())
75
76         # Check how the generator is doing by saving G's output on fixed_noise

```

```

77         if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
78             with torch.no_grad():
79                 fake = netG(fixed_noise).detach().cpu()
80                 img_list.append(vutils.make_grid(fake, padding=2, normalize=True))
81
82             iters += 1
83
84 torch.save(netD.state_dict(), 'model/netD.pth')
85 torch.save(netG.state_dict(), 'model/netG.pth')

```

executed in 2h 13m 55s, finished 00:56:18 2022-01-06

...


4. 结果查看与保存

```

In [11]: 1 plt.figure(figsize=(10,5))
2         plt.title("Generator and Discriminator Loss During Training")
3         plt.plot(G_losses, label="G")
4         plt.plot(D_losses, label="D")
5         plt.xlabel("iterations")
6         plt.ylabel("Loss")
7         plt.legend()
8         # plt.savefig("./result/Loss.jpg")
9         plt.show()

```

executed in 208ms, finished 00:56:18 2022-01-06



Visualization of G's progression


Remember how we saved the generator's output on the fixed_noise batch after every epoch of training. Now, we can visualize the training progression of G with an animation. Press the play button to start the animation.

```
In [12]: 1 %%capture
2 fig = plt.figure(figsize=(8,8))
3 plt.axis("off")
4 ims = [[plt.imshow(np.transpose(i, (1,2,0)), animated=True)] for i in img_list]
5 ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000, blit=True)
6
7 HTML(ani.to_jshtml())
```

executed in 19.7s, finished 00:56:38 2022-01-06

Animation size has reached 21398016 bytes, exceeding the limit of 20971520.0. If you're sure you want a larger animation embedded, set the animation.embed_limit rc parameter to a larger value (in MB). This and further frames will be dropped.

Out[12]:





Real Images vs. Fake Images

Finally, lets take a look at some real images and fake images side by side.

```
In [13]: 1 # Grab a batch of real images from the dataloader
2 real_batch = next(iter(dataloader))
3
4 # Plot the real images
5 plt.figure(figsize=(15, 15))
6 plt.subplot(1, 2, 1)
7 plt.axis("off")
8 plt.title("Real Images")
9 plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=5, normalize=True).cpu(),
10
11 # Plot the fake images from the last epoch
12 plt.subplot(1, 2, 2)
13 plt.axis("off")
14 plt.title("Fake Images")
15 plt.imshow(np.transpose(img_list[-1], (1, 2, 0)))
16 plt.savefig("./result/face.jpg")
17 plt.show()
```

executed in 3.19s, finished 00:56:41 2022-01-06



```
In [14]: 1 for i in range(1, len(img_list)):
2     plt.figure(figsize=(15, 15))
3     plt.axis("off")
4     plt.title("Fake Images")
5     plt.imshow(np.transpose(img_list[i], (1, 2, 0)))
6     plt.savefig(f"./result/face_{i}.jpg")
```

executed in 2m 11s, finished 00:58:52 2022-01-06

最终选出的我最满意的5张图片如下:

Fake Images



5. 计算 FID

```
In [ ]: 1 test = vutils.make_grid(fake, padding=0, normalize=True)
        2 vutils.save_image(test, './test.jpg', normalize=False)
```

In [68]:

```
1 '''
2 将一张图片切为64张图
3 '''
4 from PIL import Image
5 import sys
6
7 #将图片填充为正方形
8 def fill_image(image):
9     width, height = image.size
10    #选取长和宽中较大值作为新图片的
11    new_image_length = width if width > height else height
12    #生成新图片[白底]
13    new_image = Image.new(image.mode, (new_image_length, new_image_length), color='white')
14    #将之前的图粘贴在新图上, 居中
15    if width > height: #原图宽大于高, 则填充图片的竖直维度
16        # (x,y) 二元组表示粘贴上图相对下图的起始位置
17        new_image.paste(image, (0, int((new_image_length - height) / 2)))
18    else:
19        new_image.paste(image, (int((new_image_length - width) / 2), 0))
20
21    return new_image
22
23 #切图
24 def cut_image(image):
25     width, height = image.size
26     item_width = 64
27     box_list = []
28     # (left, upper, right, lower)
29     for i in range(0, 8): #两重循环, 生成64张图片基于原图的位置
30         for j in range(0, 8):
31             #print((i*item_width, j*item_width, (i+1)*item_width, (j+1)*item_width))
32             box = (j*item_width+k, i*item_width+k, (j+1)*item_width+k, (i+1)*item_width+k)
33             box_list.append(box)
34
35     image_list = [image.crop(box) for box in box_list]
36     return image_list
37
38 #保存
39 def save_images(image_list):
40     index = 1
41     for image in image_list:
42         image.save('./img/' + str(index) + '.png', 'PNG')
43         index += 1
44
45 file_path = "test.jpg"
46 image = Image.open(file_path)
47 # image.show()
48 image_list = cut_image(image)
49 save_images(image_list)
50
```

executed in 88ms, finished 10:57:36 2022-01-06

In [66]:

```
1 ! python fid_score.py image_data/face img --gpu 0
```

executed in 27ms, finished 10:57:23 2022-01-06

计算得到结果为 194.7550522848991

In []:

```
1
```